

Assign4_7323

February 5, 2026

```
[66]: # Dushyant Krishna Sharma
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error, r2_score
```

```
[67]: # data description
data = pd.read_csv("BostonHousing.csv")
print(data.head())
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	\
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1	296	15.3	
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2	242	17.8	
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2	242	17.8	
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3	222	18.7	
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3	222	18.7	

	B	LSTAT	MEDV
0	396.90	4.98	24.0
1	396.90	9.14	21.6
2	392.83	4.03	34.7
3	394.63	2.94	33.4
4	396.90	NaN	36.2

```
[68]: # null values
print(data.isnull().sum())
```

CRIM	20
ZN	20
INDUS	20
CHAS	20
NOX	0
RM	0
AGE	20

```
DIS          0
RAD          0
TAX          0
PTRATIO      0
B            0
LSTAT        20
MEDV         0
dtype: int64
```

```
[69]: # removing null values
data = data.dropna()
```

```
[70]: X = data.drop("MEDV", axis=1)
y = data["MEDV"]
```

```
[71]: X_train, X_test, y_train, y_test = train_test_split(
        X, y, test_size=0.2, random_state=42
)
```

```
[72]: model = LinearRegression()
model.fit(X_train, y_train)
```

```
[72]: LinearRegression()
```

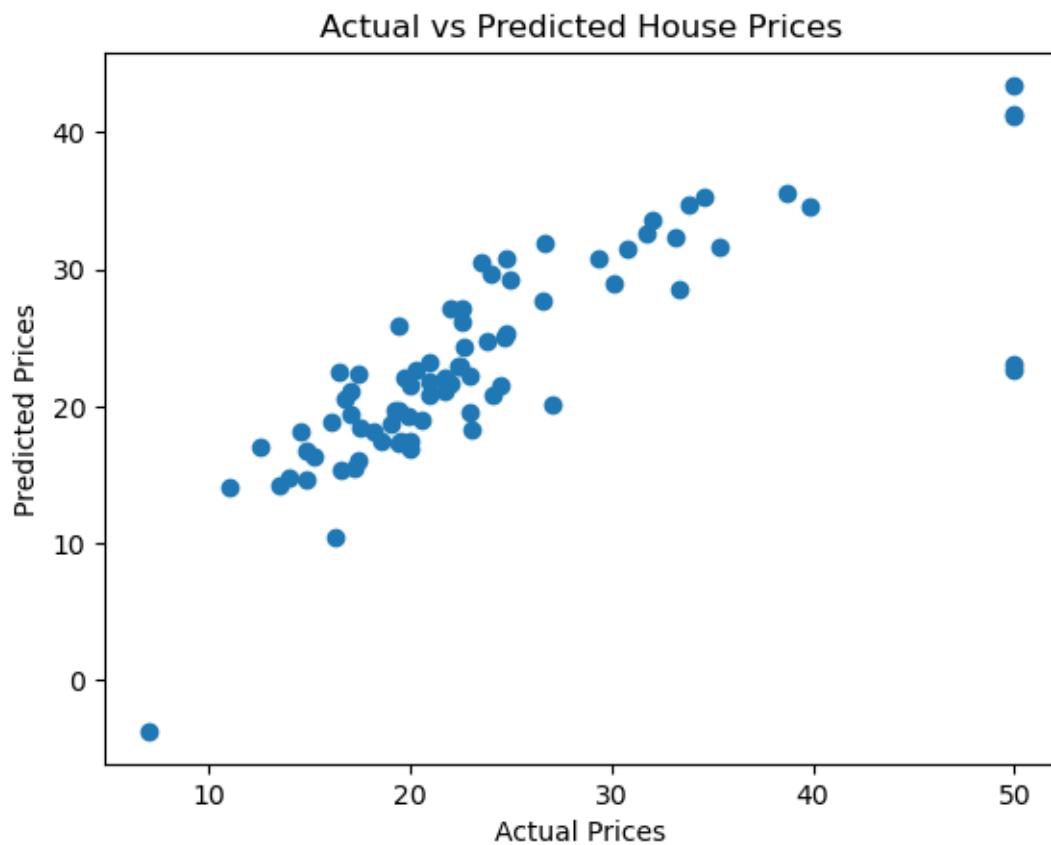
```
[73]: y_pred = model.predict(X_test)
```

```
[74]: #printing mean square , root mean square and r2 score
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
r2 = r2_score(y_test, y_pred)

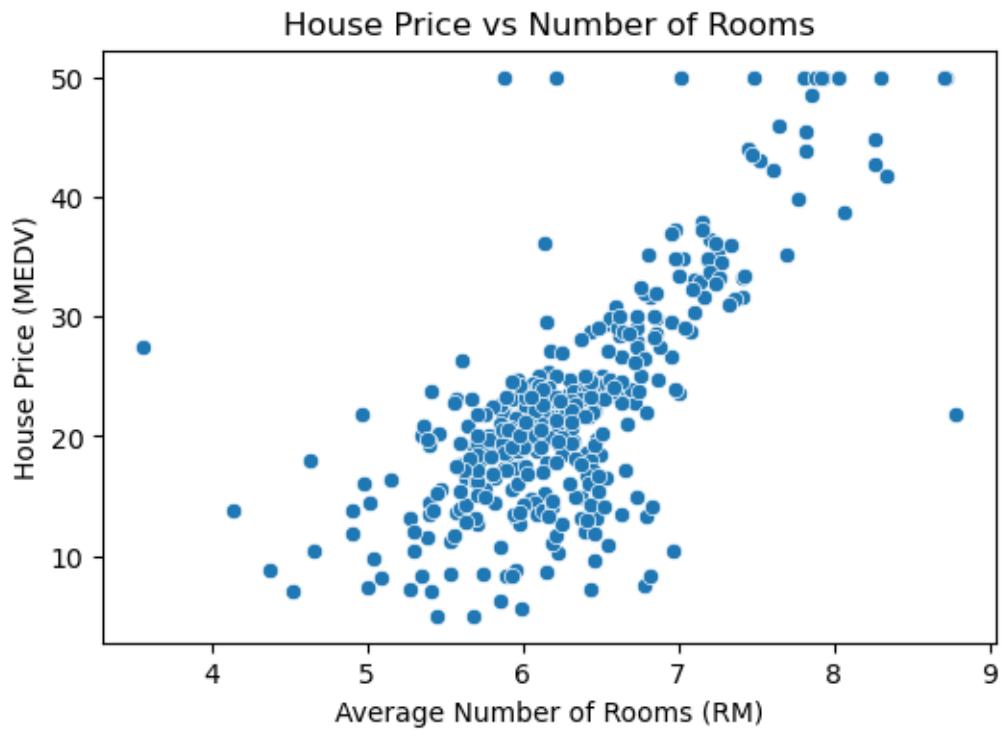
print("Mean Squared Error:", mse)
print("Root Mean Squared Error:", rmse)
print("R^2 Score:", r2)
```

```
Mean Squared Error: 31.454047664950867
Root Mean Squared Error: 5.608390826694486
R^2 Score: 0.6270849941673191
```

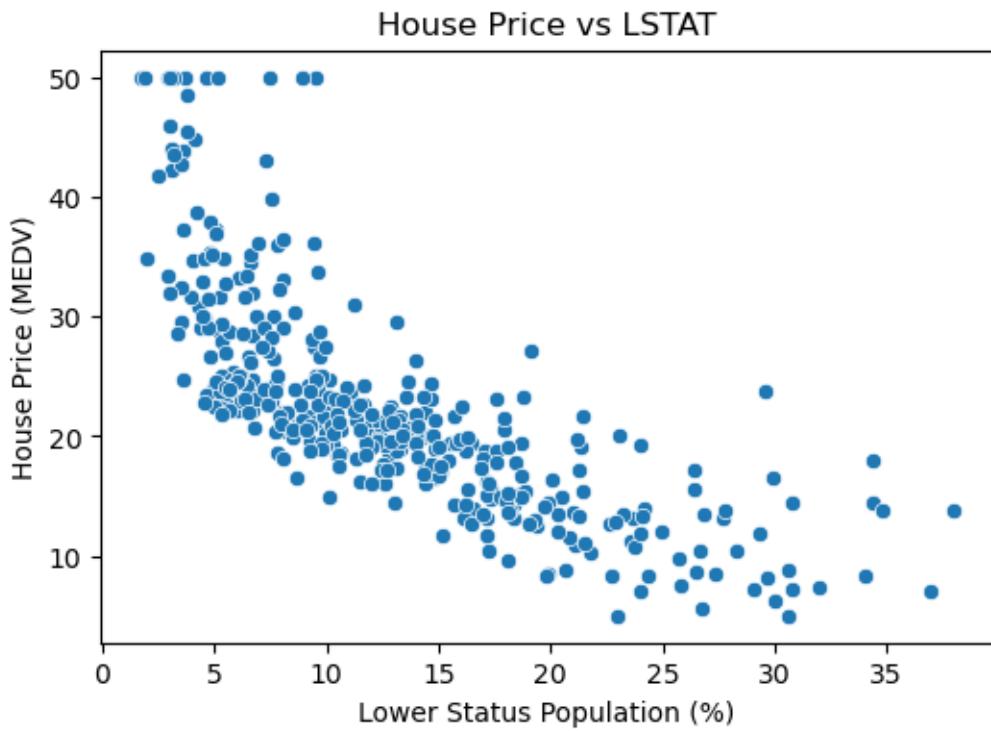
```
[75]: # actual vs predicted house prices
plt.scatter(y_test, y_pred)
plt.xlabel("Actual Prices")
plt.ylabel("Predicted Prices")
plt.title("Actual vs Predicted House Prices")
plt.show()
```



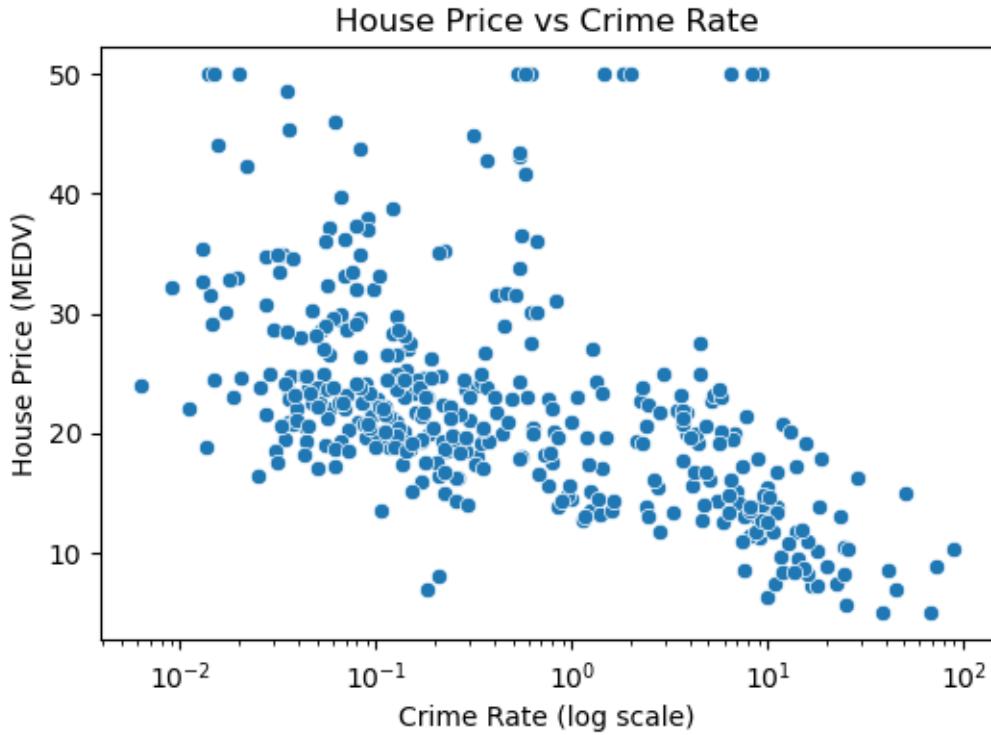
```
[76]: plt.figure(figsize=(6, 4))
sns.scatterplot(x=data["RM"], y=data["MEDV"])
plt.xlabel("Average Number of Rooms (RM)")
plt.ylabel("House Price (MEDV)")
plt.title("House Price vs Number of Rooms")
plt.show()
```



```
[77]: plt.figure(figsize=(6, 4))
sns.scatterplot(x=data["LSTAT"], y=data["MEDV"])
plt.xlabel("Lower Status Population (%)")
plt.ylabel("House Price (MEDV)")
plt.title("House Price vs LSTAT")
plt.show()
```



```
[101]: plt.figure(figsize=(6, 4))
sns.scatterplot(x=data["CRIM"], y=data["MEDV"])
plt.xscale("log")
plt.xlabel("Crime Rate (log scale)")
plt.ylabel("House Price (MEDV)")
plt.title("House Price vs Crime Rate")
plt.show()
```



```
[115]: features = data.drop("MEDV", axis=1).columns
results = []

for feature in features:
    X = data[[feature]]
    y = data["MEDV"]

    X_train, X_test, y_train, y_test = train_test_split(
        X, y, test_size=0.2, random_state=12
    )

    model = LinearRegression()
    model.fit(X_train, y_train)

    y_pred = model.predict(X_test)

    mse = mean_squared_error(y_test, y_pred)
    rmse = np.sqrt(mse)
    r2 = r2_score(y_test, y_pred)

    results.append([feature, mse, rmse, r2])
```

```
[117]: results_df = pd.DataFrame(
    results,
    columns=["Feature", "MSE", "RMSE", "R2 Score"]
)

results_df = results_df.sort_values(by="R2 Score", ascending=False)

print("\nSINGLE LINEAR REGRESSION RESULTS (ALL PARAMETERS)")
print(results_df)
```

SINGLE LINEAR REGRESSION RESULTS (ALL PARAMETERS)

	Feature	MSE	RMSE	R2 Score
5	RM	24.072058	4.906328	0.551473
12	LSTAT	24.114979	4.910700	0.550673
2	INDUS	36.548624	6.045546	0.319001
9	TAX	37.099593	6.090944	0.308735
10	PTRATIO	43.492246	6.594865	0.189623
8	RAD	44.723244	6.687544	0.166686
6	AGE	45.554442	6.749403	0.151198
0	CRIM	46.435963	6.814394	0.134773
11	B	47.426302	6.886676	0.116321
1	ZN	47.726652	6.908448	0.110724
4	NOX	47.879281	6.919486	0.107880
7	DIS	51.328924	7.164421	0.043604
3	CHAS	56.800846	7.536634	-0.058352

[]: